



redhat

JBOSS ENTERPRISE

DATA GRID

Ruben Ghio
rghio@redhat.com

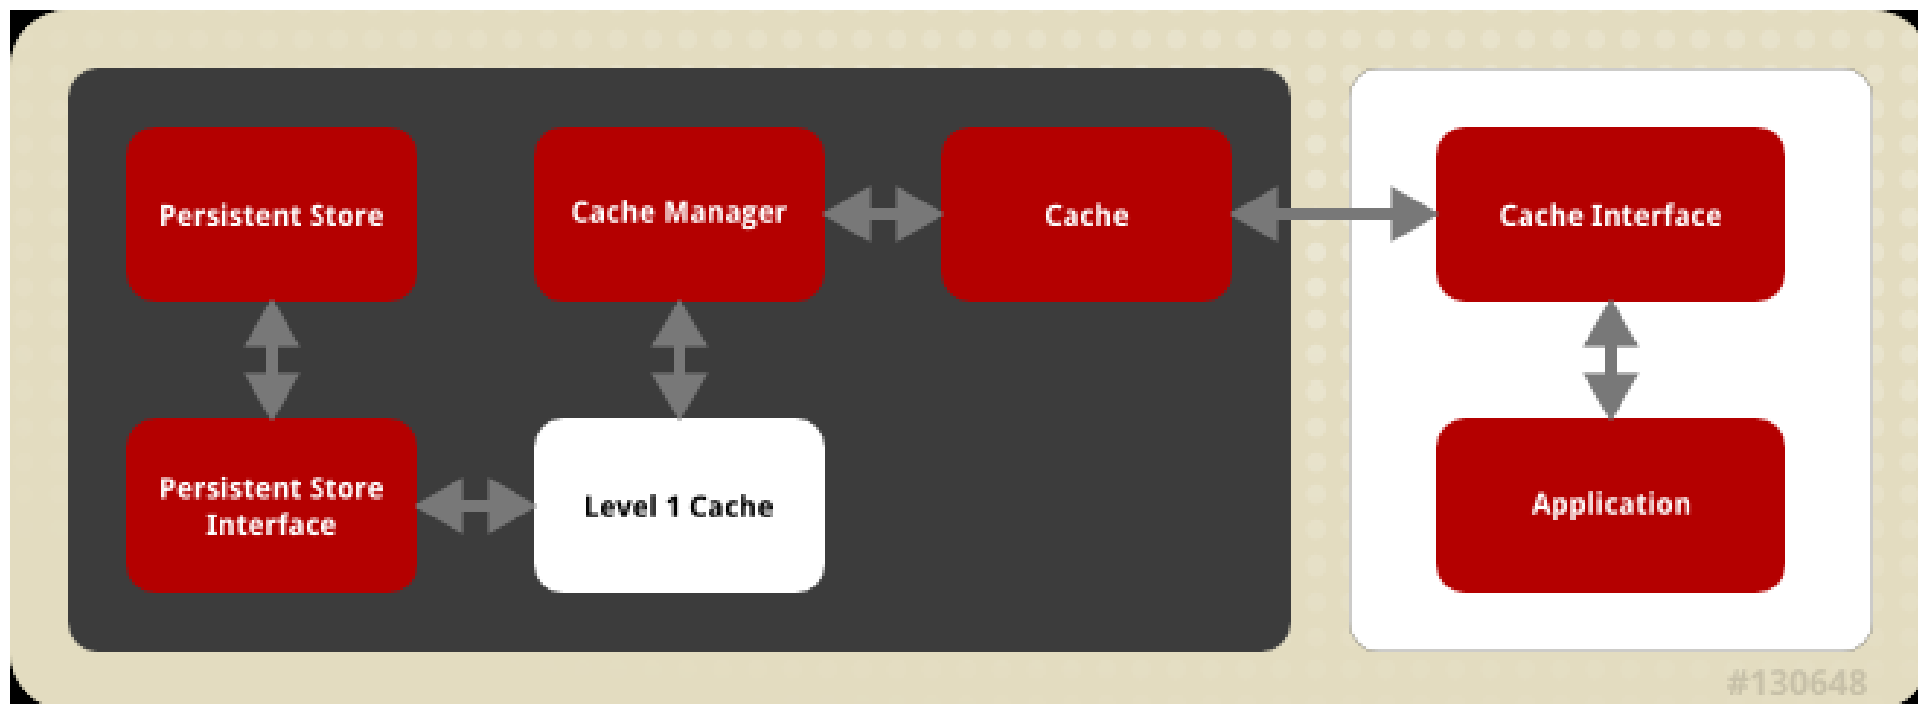
AGENDA

- WHAT ?
- HOW ?
- WHY ?
- WHEN ?

WHAT?

Highly scalable, distributed, in-memory database that lets you cost-effectively scale Big Data without making compromises. The in-memory data grid easily manages your ever-growing data stores in a high-performance, cost-effective way.

HOW?



...more HOW

Cache Interface

- Memcached
- Hot Rod
- REST

...more more HOW

Cache Manager

- Multiple Cache Managers
- Single Cache Manager

...last HOW

Persistent Store

- Cache loader
- Cache store

WHY?

- A schema-less key value store
- Grid-based data storage
- High performance
- Elastic scaling
- Flexibility and accessibility
- Transactional support
- Manageability

WHEN?

- vBETA - Currently available
- v6.0.GA - Mid June (SUMMIT ¿?)
- v6.1.GA - TBD
-

Questions?

