

Ease into noSQL,  
making big data work.

# The Challenge of scale

1) Throw “more” at the problem

2) Start Over

3) Deploy a smarter caching strategy

# The Challenge of scale

## 1) Throw “more” at the problem

More servers, more databases, etc

Expensive, may only work once

May not ever work

# The Challenge of scale

*1) Throw “more” at the problem*

## **2) Start Over**

Sharding, break up app, go stateless

May take more time than you have

Costly to rearchitect / redevelop

# The Challenge of scale

*1) Throw “more” at the problem*

*2) Start Over*

**3) Deploy a smarter caching strategy**

**NoSQL, Distributed Cache, Data Grids**

**Elastic demand based scaling**

**Highly distributed and fault tolerant**

# Smarter Caching Options

## Local Caching

How:  
Part of application

Benefits:  
Familiar to devs  
Used all the time

However,  
Single instance  
One of many other  
server tasks

## Distributed Caching

How:  
Open Source  
memcached

Benefits:  
Low / No Cost  
Massive Scale

However,  
Requires additional  
infrastructure  
May lack elasticity

## Data Grids

How:  
Standalone Products

Benefits:  
Fault tolerant  
Elastic scaling  
Managable

By 2014, at least 40% of large organizations will have deployed one or more in-memory data grids.

SOURCE – Predicts 2012 – Cloud and In Memory  
Drive Innovations in Application Platforms  
Gartner 2012

# JBoss Data Grid (JDG)

transactional data access and caching services for high-performance

on demand, automatic failover, optionally store to and load from c

nchronous or asynchronous operation. JBoss ON support for pe



# Why use JDG

**Improve scale and performance** – Move data and processing closer, partition workload across machines. JDG manages the complexity of locality, access patterns and persistence models.

**Reduce DBMS costs** – RDBMS development time, licensing, support teams don't need to grow with the application

**Simplify Development** – Application teams focus on application, JDG manages persistence in an appropriate fashion.

# Existing Uses

Delivery of **real time** routing, tracking and logistics **information**

Provide rapid access to content for **thousands of users**

Reduce RDBMS workload with **intelligent caching**

Enable **in-memory speed** with **transactional support** for compute intensive applications

**TBD – Add CBOE use-case**

***Okay, I need that in captain dummy talk, Kaylee***

**In-JVM, self-managed cache to scale applications otherwise bound  
by a relational database**

***Pacemaker for a database***

**Distributed Key/Value store to support operational workloads using a  
standard, language-agnostic API**

***a NoSQL data store***

# JDG Features

Configurable data distribution

Overflow memory to disk or other 'cachestore'

Integrates with existing data stores

Multiple locking strategies

Scale linearly, without complexity

Low latency - **Fast and Deterministic** access

Highly Available and resilient


Elastic – scale out and back again

Manageable – JBoss ON is JDG aware

# Using JBoss Data Grid

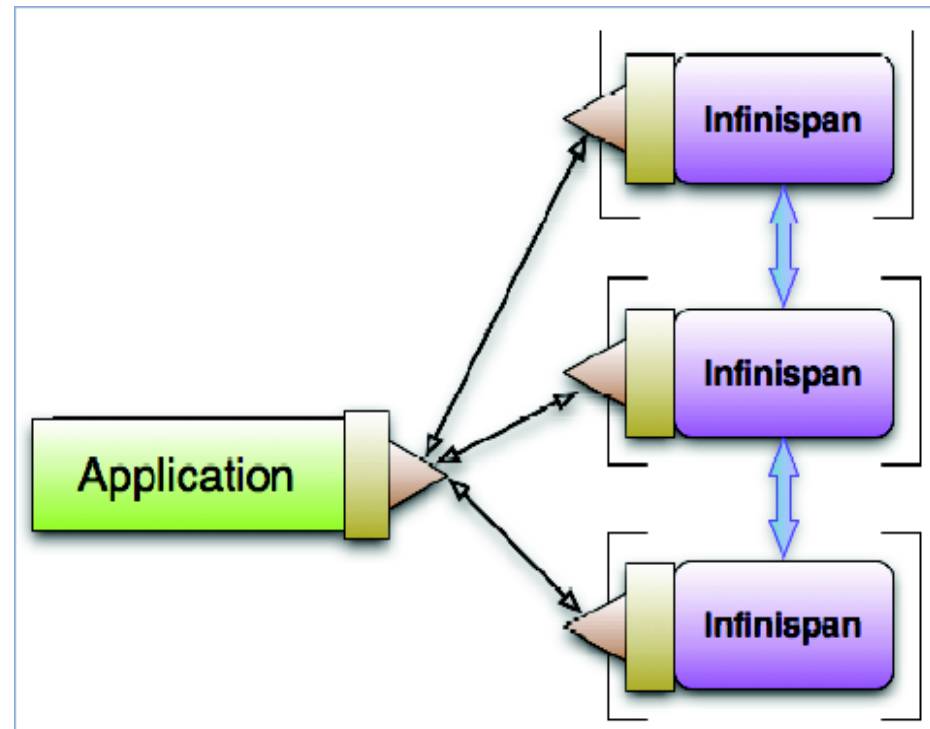
## Option 1 – As a library

Application co-located with cache node



## Option 2 – Remote Client/Server

Remote Client Access, multiple protocols



# Library Features - In-memory data grid

## Distributed, Off-Heap HashMap

Application container agnostic

InVM - Java Only

Synchronous/asynchronous API

Event aware

Data can be grouped

JTA compatible

JMX monitoring

# Library Benefits

## Performance booster

In-memory access for data that is expensive to retrieve and frequently accessed

## Simplify Large Scale Applications


Built-in eviction to manage JVM heap

## Customizable


Very rich API - data grouping, notifications, serialization

## Supporting Scaffolding for existing frameworks

CDI, SEAM, Apache Camel, ESBs



# Remote Client / Standalone Grid Features



**Client/Server topology**

**REST/Memcache/Hot Rod client access**

**Manageable via JON**




# Standalone Grid Benefits

**Fast** - In-memory, low latency, high concurrency

**Available** - Inherent redundancy

**Distributed** - Data locality

Data managed independently of application



# Available Grid APIs

	Protocol	Client Libraries	Clustered ?	Smart Routing	Load Balancing/ Failover
<b>REST</b>	<b>Text</b>	N/A	<b>Yes</b>	<b>No</b>	Any HTTP load balancer
<b>Memcached</b>	<b>Text</b>	Plenty	<b>Yes</b>	<b>No</b>	Only with predefined server list
<b>Hot Rod</b>	<b>Binary</b>	Currently only Java	<b>Yes</b>	<b>Yes</b>	Dynamic

Supported Beta – out **now**

GA – June (before Summit)

6.x – Summery

Library mode out of Tech Preview

6.x + 1 – Q4 2012

# JDG Beta

# OUT NOW

